

Visual Odometry Pipeline

Pascal Buholzer, Fabio Dubois, Milan Schilling, Miro Voellmy

January 8, 2017

Symbols

@Miro - we can also delete this chapter

Contents

1	Introduction	4
2	Implementation	5
2.0.1	Coordinate Frames	5
2.0.2	Pipeline overview	5
2.0.3	Conventions	5
2.0.4	Options and parameters	6
2.1	Initialization	7
2.1.1	Bootstrapping	7
2.1.2	findCorrespondences()	7
2.1.3	eightPointRansac	7
2.1.4	linearTriangulation	8
2.1.5	bundleAdjust	8
2.1.6	applySphericalFilter	8
2.2	Continuous Operation	10
2.2.1	Relative pose estimation	10
2.2.2	find_correspondences_cont()	10
2.2.3	updateKpTracks()	10
2.2.4	triangulateNewLandmarks()	10
2.2.5	Bundle Adjustment	10
2.2.6	Reinitialization	11
3	Results	12
3.1	Bootstrapping methods	12
3.2	Overall performance	12
4	Discussion	13
4.0.1	Correspondence search	13
4.0.2	Bundle Adjust	13
4.0.3	Pose estimation algorithm	13
4.0.4	Reinitialization	13
4.0.5	Feature work	13
5	Conclusion	13

1 Introduction

@Miro @Milan, read it trough please

During this mini project a monocular visual odometry pipeline was developed. This pipeline takes the consecutive gray-scale images of a single digital camera as input. The output of the pipeline is the position of the camera in relation to its initial position for each frame. The pipeline is programmed in such a way that the Markov assumption is valid. This means that the current computation step is only dependent on the previous step to reduce the required computation effort.

keywords: (VO, sequential,monocular, markov assumption)

2 Implementation

This pipeline was developed in MATLAB. Since the group consisted of four students, a Git repository was used to be able to work on different files simultaneously, and to enable version control.

2.0.1 Coordinate Frames

In this mini project the coordinate frames were defined as shown in ???. The camera coordinates are in a way oriented, that the x-y plane lies parallel to the image plane, while the z-axis is pointing towards the scenery. The world frame however is oriented in such a way that the x-y plane is parallel to the ground and the z-axis is pointing upwards. The origin of the world frame is at the same location as the origin of the first bootstrap image.

Transformation between frames are described by homogenous transformation matrices. T_{AB} maps points from frame B to frame A .




Figure 1: Coordinate Frames

2.0.2 Pipeline overview

As shown in ?? the pipeline consists mainly of three parts:

1. Bootstrapping (see ??)
2. Initialization (see ??)
3. Continuous operation (see ??)

@Milan is this flowchart up to date? If yes all good. :-)

2.0.3 Conventions

- Index of previous frame: i
- Index of current frame: j
- Index of frame for newly added candidate keypoint: *first*
- Pose difference between previous to current frame: $T_{C_i C_j}$
- $[u/v]$: Pixel coordinates
- Query keypoints: Keypoints newly generated in frame j
- Candidate keypoint: A keypoint without associated landmark
- Harris Matcher: Descriptor matching keypoint tracker (based on Harris features) developed during the lecture.




Figure 2: Overview flow chart

2.0.4 Options and parameters

The pipeline was designed in a modular way. Key algorithms are abstracted into self-contained functions, as described in the pipeline overview. Next to this 'functional programming' approach all the tuning variables (e.g. number of keypoints, bearing angle thresholds, etc.) were centrally aggregated in a parameter struct. For further insight please consult the function `loadParameters.m`.

In order to run the visual odometry two launch procedures were implemented:

- **Debug Mode:** For development and debug mode the `main.m` with default parameters can be executed. Numerous individual plots are displayed with insightful information about matching, inlier rejection and triangulation.
- **Simple GUI:** Out of performance reasons a more compact and user-friendly display of the pipeline output was created with a GUI designed with the Matlab GUIDE application, see figure ?? . Only the most crucial entities, like number of landmarks, are visualized for intuitive understanding.

How to run the GUI Please follow the steps below to run the visual odometry through the GUI environment:

1. adapt dataset paths in `loadParameters.m`
2. type into MATLAB command window: `gui_simple`
3. in *Parameters* panel select dataset to run on and toggle respective radio buttons
4. hit *Run* to trigger the visual odometry

Advanced parameter tuning can be achieved by changing the default parameters in `loadParameters.m`.




Figure 3: Graphical user interface
all features (red ●), inlier features (green ●), triangulated features (green ×)

2.1 Initialization

The first step of the initialization is the bootstrapping which outputs an imagepair for further processing. This imagepair varies depending on the dataset and is required since the baseline between consecutive images is often too small for accurate landmark triangulation and pose estimation.

Features are then matched across the two bootstrap images and the pose of the second camera is estimated using 8-point-Ransac. Landmarks are generated using the pose and the matched features. These landmarks and the pose are then refined using bundle adjustment. To end the initialization the landmarks which are unrealistic are discarded.

2.1.1 Bootstrapping

In order to pick reasonable image pairs for initialization an automatic procedure was implemented. Since a good initialization relies heavily on the number of inlier keypoints a hard constrain on their minimum number of `min_num_inlier_kp = 600` was set.

Two approaches were investigated:

- *SSD matching + Baseline/Depth*:
correspondence search via sum of squared differences (SSD) on Harris features & 8-point-Ransac for outlier rejection & baseline/depth-ratio (as proposed in the lecture)

$$\frac{C1_{baseline}}{C1_{av_depth}} \geq min_b2dratio = 0.1 \quad (1)$$

- *KLT + bearing angle*:
correspondence search and inlier selection via Kanade-Lukas-Tomasi (KLT) tracking on query Harris corners & minimum average bearing angle

$$av_bearing_angle_deg \geq min_av_angle_deg = 10deg \quad (2)$$

As elaborated in ?? the second approach outperformed the other and is selected as auto-bootstrapping method.

2.1.2 findCorrespondences()

The same approaches discussed in ?? are used to find the final matching between the bootstrap frames.

2.1.3 eightPointRansac

Normalized 8-point-Ransac is used to estimate the essential matrix and filter the outliers of the keypoint matching. The error function consists of the distance from the query keypoint to the epipolar line. All keypoints which have an error smaller than a certain threshold are counted as inliers. The final inlier-set is the one with the most inliers. The essential matrix obtained by the 8-point-Ransac is then decomposed with respect to the obtained translation such that we get the final transformation.

2.1.4 linearTriangulation

3D-landmarks are generated using linear triangulation of the previously matched keypoints and transformation.

2.1.5 bundleAdjust

The landmarks and transformation between the first to frames is optimized using bundle adjustment.

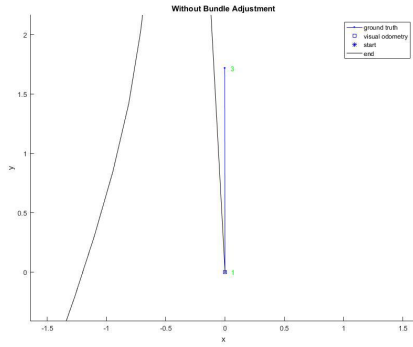


Figure 4: Init Flow chart

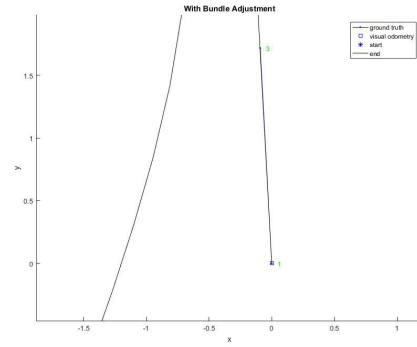


Figure 5: Init Flow chart

2.1.6 applySphericalFilter

Discards all landmarks which are not within a half-sphere in front of the camera. This removes on one hand landmarks triangulated behind the camera (negative z-components) and also landmarks further away than the filter cutoff radius. With this filtering only landmarks in the neighborhood are kept, which are more reliably triangulated (smaller uncertainty cone). Note, that this absolute thresholding relies on an accurate scale estimation.

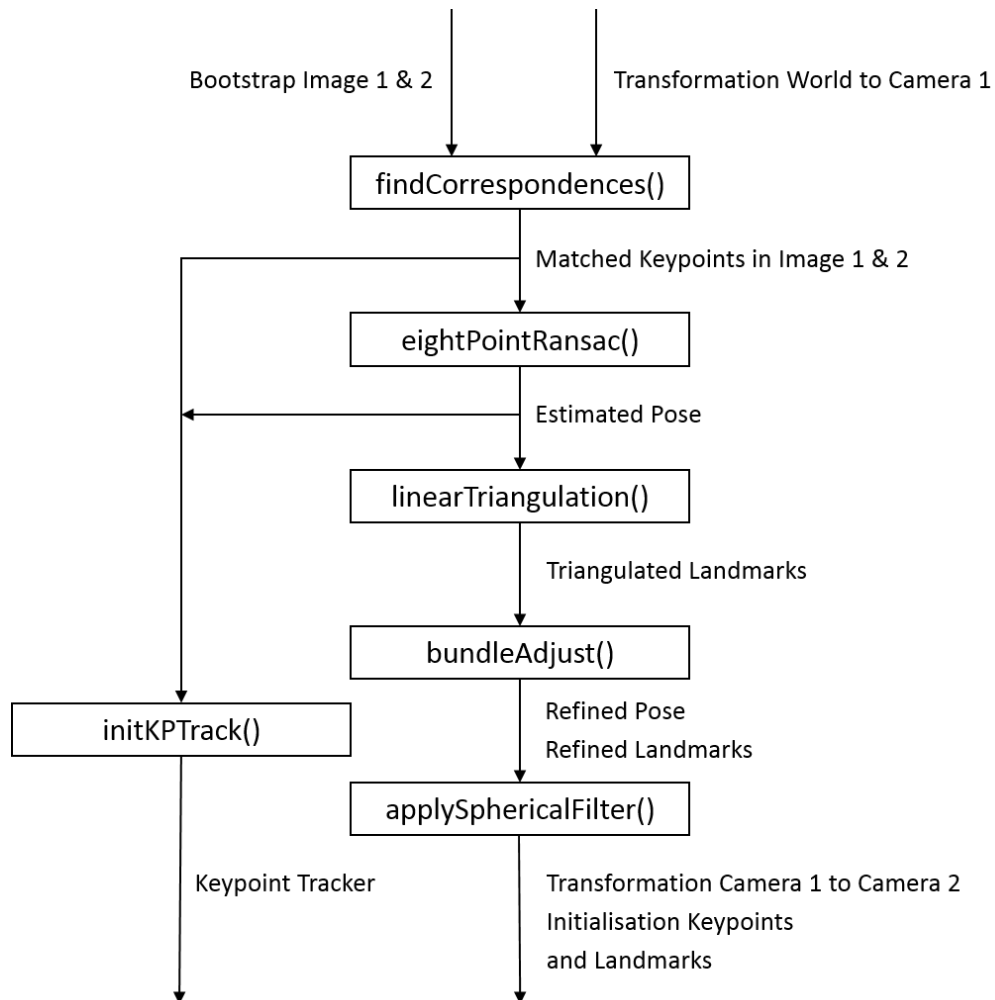


Figure 6: Init Flow chart

2.2 Continuous Operation

Continuous operation of the VO pipeline is implemented in the 'process frame' function. It tracks keypoints with corresponding landmarks over several frames while estimating the pose difference between successive frames. Further, a keypoint tracker finds new candidate keypoints which will become new landmarks if a candidate keypoint was tracked far enough and achieved 'good' triangulability. This ensures to never run out of landmarks and keypoints if the image changes over time. The routines of the continuous operation shown in ?? are described below.

2.2.1 Relative pose estimation

To estimate the pose difference $T_{C_i C_j}$ from frame i to j we use the p3p-RANSAC algorithm also used in exercise 5. If wished the RANSAC can also use DLT pose estimation. Using these RANSAC algorithm ensures to remove outliers from our landmarks. We don't use DLT refinement after the p3p RANSAC since the best guess from p3p often gave better results.

2.2.2 `find_correspondences_cont()`

Tracking keypoints with existing landmarks from frame i to frame j is achieved by the function *find_correspondences_cont*. The user can choose whether to use a KLT or Harris matcher. As a by-product, the generated query keypoints are saved to be used by the candidate keypoint tracker in a successive step so they don't have to be generated twice which saves computation time. The number of generated query keypoints is adjustable by a parameter. In case the KLT tracker is active (which does not return query keypoints by default) new query keypoints are generated using Harris features. The amount of newly generated keypoints is the difference of remaining and wished candidate keypoints.

2.2.3 `updateKpTracks()`

In every frame, candidate keypoints from previous frames are tracked to j -frame. Every candidate keypoint track consists of the following entries: $\{[u/v]_j, [u/v]_{first}, T_{WC_{first}}, nr_trackings\}$. After successive tracking, $[u/v]_j$ as well as the number of successful successive trackings are updated. The user can choose whether to use a KLT or a Harris matcher. In case a candidate keypoint could not be tracked it's whole track gets removed from the tracker. If there are less candidate keypoints in the tracker then desired, newly generated keypoints (generated in `find_correspondences_cont`) are added to the tracker. Every newly added keypoint is stored together with the current pose $T_{WC_{first}}$.

2.2.4 `triangulateNewLandmarks()`

In order to localize correctly (see ??), a sufficient number of landmarks is required. If the number of landmarks in the current frame drops below a certain threshold, the 'triangulateNewLandmarks' function is called to generate new landmarks from the candidate keypoint tracks. To check, whether a candidate keypoint is ready for triangulation the bearing angle between it's first and it's current observation is calculated. This is done for every keypoint candidate. If the bearing angle of a candidate keypoint exceeds a certain threshold, it is removed from the keypoint tracker and a new landmark gets triangulated with the algorithm developed in exercise 4. The bearing angle threshold is adaptive to the number of remaining landmarks. The higher the number of remaining landmarks the higher the threshold. This ensures generation of new landmarks in case the filter starts to run out of landmarks but improves triangulation results once enough landmarks are in the pipeline. To filter outliers within the newly generated landmarks (e.g. from image noise) two checks are implemented:

1. Spherical filter: Already described in ??.
2. Reprojected error filter: Landmarks passed the spherical filter get reprojected into the current frame. If the reprojected point is too far apart from the candidate keypoint coordinates it is discarded.

Landmarks and their corresponding keypoints which passed these two filters are appended to the existing landmarks and keypoints vectors and will be used during the next localization iteration.

2.2.5 Bundle Adjustment

Todo: @Miro Explain how it works - add to flowchart

2.2.6 Reinitialization

@Milan Explain how it works, and when it's needed

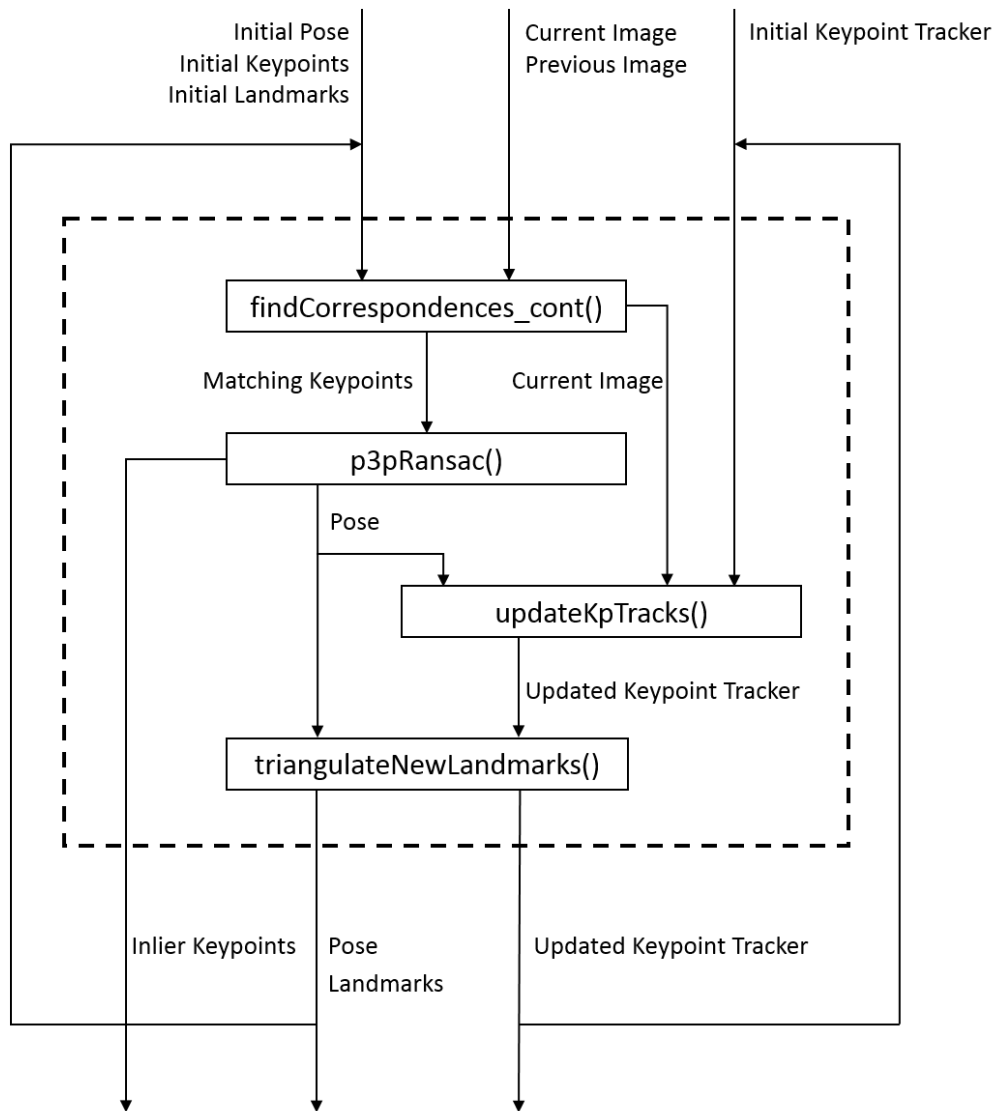


Figure 7: Cont Flow chart

3 Results

3.1 Bootstrapping methods

The comparison showed a significantly higher performance of the KLT tracker approach, this being orders of magnitude faster and yielding many more feature correspondences also over large distances. ?? depicts the index tuples retrieved through the second approach with `min_num_inlier_kp = 600` being required.

approach		Kitti	Malaga	Parking
<i>KLT + Bearing angle</i>	first idx	1	1	1
	second idx	4	6	5
	# keypoints	600	600	600
	angle [deg]	1.37	4.59	4.07

Table 1: Bootstrapping pair indices for different datasets with minimum keypoint inliers

When relaxing the minimum number of inlier keypoints allowed to reach the desired baseline/depth ratio of 10 % for the first and a bearing angle of 10 degrees for the second bootstrapping approach results are shown in ??.

approach		Kitti	Malaga	Parking
<i>SSD Harris desc. + baseline/depth ratio</i>	first idx	1	1	1
	second idx	6	10	7
	# keypoints	34	39	54
	ratio [%]	10.1	10.9	11.5
<i>KLT + bearing angle</i>	first idx	-	1	1
	second idx	-	16	17
	# keypoints	-	99	74
	angle [deg]	-	10.4	10.55

Table 2: Bootstrapping pair indices for different datasets and no keypoint number constraint

We observed that the first approach yields far too few keypoints, due to the `matchDescriptor()` method. Furthermore, the pure lateral camera motion of dataset Parking suits well this baseline/depth bootstrapping approach, which confirms our intuition given the 'triangulability condition' described in ??.

Deploying the second approach on the Kitti dataset evidently fails due to the straight camera motion, inhibiting angles ($[0.5, 3.7]$ deg) to reach the desired 10 degrees. As soon as there is a distinctive rotational movement involved, e.g. in Malaga and Parking, a useful bootstrapping is performed.

3.2 Overall performance

(keywords: Real time ness, comparison to groundtruth, compare different datasets Impact of features) Here we describe it such as it runs best in our opinion.

- Chosen key parameters
 - Kitti @Pascal
 - Malaga @Milan
 - Parking @Pascal
 - Poly-up @ Fabio
 - Poly-down @Fabio
- Runtime characteristics (e.g. average nr inlier landmarks, and new landmarks, candidate kp)
- plot of groundtruth
- 3d landmarks if possible
- speed & realtimeness, slowest functions - why?
- speed & realtimeness, slowest functions - why?

4 Discussion

4.0.1 Correspondence search

KLT vs Harris matching @ Milan

4.0.2 Bundle Adjust

improvements reached @Miro

4.0.3 Pose estimation algorithm

DLT refinement vs p3p guess in continuous maybe not @Pascal

4.0.4 Reinitialization

@Milan How often does it happen? Is the result better afterwards?

4.0.5 Feature work

A list of potential improvements @ Milan

Possible future work, improvements (loop closure, ...)

- bootstrapping with both information rotation through bearing angle diff and translation through baseline/depth ratio
- ...

5 Conclusion

What have we learned, what worked? @Milan